

```
#####
#####
# Creator: Ovando Carter - using "Neural Networks From Scratch by Harrison
Kingsley"
# Including:
# Modeling neurons and building layers *
# using dot product of a layer of neurons *
# Use of generated data *
# Use of test data *
# including activation function - ReLU *
# - Softmax
# Use of Catagorical Crossentropy Loss functions
# Optimisation - adjust the weights and biases to decrease the loss. *
# Back Propigation using chain rule *
# Optimisers - Stochastic Gradient Descent (SGD) -> using momentums
# - Adaptive gradient (AdaGrad)
# - Root Mean Square Propagation (RMSProp)
# - Adaptive Momentum (Adam) *
# Dropout
# Binary Crossentropy *
# The two main methods for calculating error in regression are
# - mean squared error (MSE)
# - and mean absolute error (MAE).
# Imports and shuffles data from directories using:
# - test data
# - training data
# Model Evaluation
# Predictions
```

```
#####
#####
# used to monitor data output
```

```
from itertools import count
import pandas as pd
import matplotlib.pyplot as plt
from matplotlib.animation import FuncAnimation
#####
#####
import nnfs
import os
```

```
import pickle # allows us to save the parameters of the model into a file.
import copy
#from nnfs.datasets import spiral_data, vertical_data, sine_data #imported data
set from nnfs.datasets
```

```
try:
    import numpy as np
except:
    # Install numpy if they do not have it installed already
    import pip
    pip.main(['install', 'numpy'])
    import numpy as np
```

```
try:
    import cv2
except:
    # Installs opencv if you do not have it installed already,
    # NB: opencv download can take 30 - 60 mins.
    import pip
    pip.main(['install', 'opencv-python'])
    import cv2
```

```
nnfs.init()
```

```
#####
#####
# Initialise weights
#####
#####
```

```
'''
```

Keep values small by normalising them

randn is a gaussian distribution bounded around 0. We multiplied 0.10 since some of the values that came out were greater than one. Yet we were using a gaussian distribution bounded around zero.

```
'''
```

```
class layer_Dense:
    def __init__(self, n_inputs, n_neurons,
                  weight_regularizer_l1 = 0, weight_regularizer_l2 = 0,
```

```

        bias_regularizer_l1 = 0, bias_regularizer_l2 = 0):
self.weights = 0.01 * np.random.randn(n_inputs, n_neurons)
self.biases = np.zeros((1, n_neurons))
# Initialize weights and biases
self.weight_regularizer_l1 = weight_regularizer_l1
self.weight_regularizer_l2 = weight_regularizer_l2
self.bias_regularizer_l1 = bias_regularizer_l1
self.bias_regularizer_l2 = bias_regularizer_l2

# Forward pass
def forward(self, inputs, training):
    self.output = np.dot(inputs, self.weights) + self.biases
    self.inputs = inputs # added based on the book. supposed to help with
back propagation

# Backward pass - backpropagation - derivatives
def backward(self, dvalues):
    # Gradients on parameters
    self.dweights = np.dot(self.inputs.T, dvalues)
    self.dbiases = np.sum(dvalues, axis = 0, keepdims = True)

    # Gradients on regularizations
    # L1 on weights
    if self.weight_regularizer_l1 > 0:
        dl1 = np.ones_like(self.weights)
        dl1[self.weights < 0] = -1
        self.dweights += self.weight_regularizer_l1 * dl1
    #L2 on weights
    if self.weight_regularizer_l2 > 0:
        self.dweights += 2 * self.weight_regularizer_l2 * \
            self.weights
    # L1 on biases
    if self.bias_regularizer_l1 > 0:
        dl1 = np.ones_like(self.biases)
        dl1[self.biases < 0] = -1
        self.dbiases += self.bias_regularizer_l1 * dl1
    # L2 on biases
    if self.bias_regularizer_l2 > 0:
        self.dbiases += 2 * self.bias_regularizer_l2 * \
            self.biases

    # Gradient on values
    self.dinputs = np.dot(dvalues, self.weights.T)

# Retrive layer parameters
def get_parameters(self):

```

```

        return self.weights, self.biases

# Set weights and biases in a layer instance
def set_parameters(self, weights, biases):
    self.weights = weights
    self.biases = biases

#####
#####
# Layer Dropout
#####
#####

# Dropout
class Layer_Dropout:

    #Init
    def __init__(self, rate):
        # Store rate, we invert it as for example dropout
        # of 0.1 we need success rate of 0.9
        self.rate = 1 - rate

    # Forward pass
    def forward(self, inputs, training):
        # Save input values
        self.inputs = inputs

        # If not in the training mode - return values
        if not training:
            self.output = inputs.copy()
            return

        #Generate and save scaled mask
        self.binary_mask = np.random.binomial(1, self.rate,
                                                size = inputs.shape)/self.rate
        # Apply mask to output values
        self.output = inputs * self.binary_mask

    # Backward pass
    def backward(self, dvalues):
        # Gradient on values
        self.dinputs = dvalues * self.binary_mask

# Input "layer"
class Layer_Input:

```

```

#Forward Pass
def forward(self, inputs, training):
    self.output = inputs

#####
#####
# Activation functions
#####
#####
# Rectified Linear Activation Function
class Activation_ReLU:
    # Forward pass
    def forward(self, inputs, training):
        self.inputs = inputs # Remember input values - added to help with back
propagation
        self.output = np.maximum(0, inputs)
    # Backward pass
    def backward(self, dvalues):
        # Since we need to modify the original variable,
        # let's make a copy of the values first
        self.dinputs = dvalues.copy()

        # Zero gradient where input values were negative
        self.dinputs[self.inputs <= 0] = 0

    # Calculate predictions for outputs
    def predictions(self, outputs):
        return outputs

# Exponential activation function
class Activation_Softmax:
    def forward(self, inputs, training):
        # Remember input values
        self.inputs = inputs
        # Get unnormalized probabilities (axis = 1, keepdims=True used for
normalisations)
        exp_values = np.exp(inputs - np.max(inputs, axis = 1, keepdims=True))
        # Normalize them for each sample
        probabilities = exp_values / np.sum(exp_values, axis = 1, keepdims=True)
        self.output = probabilities

```

```

# Backwards pass
def backward(self, dvalues):

    # Create uninitialized array
    self.dinputs = np.empty_like(dvalues)

    # Enumerate outputs and gradients
    for index, (single_output, single_dvalues) in \
        enumerate(zip(self.output, dvalues)):
        # Flatten output array
        single_output = single_output.reshape(-1, 1)
        #Calculate Jacobian matrix of the output and
        jacobian_matrix = np.diagflat(single_output) - \
            np.dot(single_output, single_output.T)
        #Calculate sample-wise gradient
        # and add it to the array of sample gradients
        self.dinputs[index] = np.dot(jacobian_matrix, single_dvalues)

# Calculate predictions for outputs
def predictions(self, outputs):
    return np.argmax(outputs, axis=1)

# Sigmoid activation
class Activation_Sigmoid:

    # Forward pass
    def forward(self, inputs, training):
        # Save input and calculate/save output
        # of the sigmoid function
        self.inputs = inputs
        self.output = 1 / (1+ np.exp(-inputs))

    # Backward pass
    def backward(self, dvalues):
        # Derivative - calculates from output of the sigmoid function
        self.dinputs = dvalues * (1 - self.output) * self.output

    # Calculate predictions for outputs
    def predictions(self, outputs):
        return (outputs > 0.5) * 1

# Linear activation
class Activation_Linear:

    # Forward pass
    def forward(self, inputs, training):

```

```

    # Just remember values
    self.inputs = inputs
    self.output = inputs

    # backward pass
    def backward(self, dvalues):
        # derivative is 1, 1 * dvalues = dvalues - the chain rule
        self.dinputs = dvalues.copy()

    # Calculate predictions for outputs
    def predictions(self, outputs):
        return outputs

#####
#####
# Stochastic Gradient Descent (SGD)
#####
#####

class Optimizer_SGD:
    # Initialize optimizer - set settings,
    # learning rate of 1. is default for this optimizer - changing learner rate can
    effect the ability for the
    # model to find global minimums and not get stuck in any local minimums.
    The aim is to get a low loss and a higher accuracy.
    def __init__(self, learning_rate = 1.0 , decay = 0., momentum = 0.):
        self.learning_rate = learning_rate
        self.current_learning_rate = learning_rate
        self.decay = decay
        self.iterations = 0
        self.momentum = momentum

    # Call once before and parameter updates - this part should reduce the
    learning rate with time
    def pre_update_params(self):
        if self.decay:
            self.current_learning_rate = self.learning_rate * \
                (1. / (1. + self.decay * self.iterations))

    # Update parameters
    def update_params(self, layer):
        # If we use momentum
        if self.momentum:
            # If layer does not contain momentum arrays, create them
            # filled with zeros

```

```

if not hasattr(layer, 'weight_momentums'):
    layer.weight_momentums = np.zeros_like(layer.weights)
    # If there is no momentum array for weights
    # The array doesn't exist for biases yet either.
    layer.bias_momentums = np.zeros_like(layer.biases)

# Build weight updates with momentum - take previous
# updates multiplied by retain factor and update with
# current gradients
weight_updates = \
    self.momentum * layer.weight_momentums - \
    self.current_learning_rate * layer.dweights
layer.weight_momentums = weight_updates

# Build bias updates
bias_updates = \
    self.momentum * layer.bias_momentums - \
    self.current_learning_rate * layer.dbiases
layer.bias_momentums = bias_updates

# Vanilla SGD updates (as before momentum update)
else:
    weight_updates = -self.current_learning_rate * \
        layer.dweights
    bias_updates = -self.current_learning_rate * \
        layer.dbiases
# Update weights and biases using either
#vanilla or momentum updates
#layer.weights += -self.learning_rate * layer.dweights
#layer.biases += - self.learning_rate * layer.dbiases
layer.weights += weight_updates
layer.biases += bias_updates

# Call once after parameter updates
def post_update_params(self):
    self.iterations += 1

#####
#####
# Adaptive gradient (AdaGrad)
#####
#####
class Optimizer_Adagrad:
    # Initialize optimizer - set settings,
    # learning rate of 1. is default for this optimizer - changing learner rate can
    effect the ability for the

```


model to find global minimums and not get stuck in any local minimums.
The aim is to get a low loss and a higher accuracy.

```
def __init__(self, learning_rate = 1. , decay = 0., epsilon = 1e-7):  
    self.learning_rate = learning_rate  
    self.current_learning_rate = learning_rate  
    self.decay = decay  
    self.iterations = 0  
    self.epsilon = epsilon
```

Call once before and parameter updates - this part should reduce the learning rate with time

```
def pre_update_params(self):  
    if self.decay:  
        self.current_learning_rate = self.learning_rate * \  
            (1. / (1. + self.decay * self.iterations))
```

Update parameters

```
def update_params(self, layer):
```

```
    # If layer does not contain cache arrays,  
    # create them filled with zeros  
    if not hasattr(layer, 'weight_cache'):  
        layer.weight_cache = np.zeros_like(layer.weights)  
        layer.bias_cache = np.zeros_like(layer.biases)
```

```
    # Update cache with squared current gradients  
    layer.weight_cache += layer.dweights**2  
    layer.bias_cache += layer.dbiases**2
```

```
    # Vanilla SGD parameter update + normalization  
    # with square rooted cache
```

```
    layer.weights += -self.current_learning_rate * \  
        layer.dweights / \  
        (np.sqrt(layer.weight_cache) + self.epsilon)  
    layer.biases += -self.current_learning_rate * \  
        layer.dbiases / \  
        (np.sqrt(layer.bias_cache) + self.epsilon)
```

Call once after parameter updates

```
def post_update_params(self):  
    self.iterations += 1
```

```
#####  
#####  
# Root Mean Square Propagation (RMSProp)  
#####
```

```
#####
class Optimizer_RMSprop:
    # Initialize optimizer - set settings,
    # learning rate of 1. is default for this optimizer - changing learner rate can
    # effect the ability for the
    # model to find global minimums and not get stuck in any local minimums.
    The aim is to get a low loss and a higher accuracy.
    def __init__(self, learning_rate = 0.001, decay = 0., epsilon = 1e-7, rho = 0.9):
        self.learning_rate = learning_rate
        self.current_learning_rate = learning_rate
        self.decay = decay
        self.iterations = 0
        self.epsilon = epsilon
        self.rho = rho

    # Call once before and parameter updates - this part should reduce the
    learning rate with time
    def pre_update_params(self):
        if self.decay:
            self.current_learning_rate = self.learning_rate * \
                (1. / (1. + self.decay * self.iterations))

    # Update parameters
    def update_params(self, layer):

        # If layer does not contain cache arrays,
        # create them filled with zeros
        if not hasattr(layer, 'weight_cache'):
            layer.weight_cache = np.zeros_like(layer.weights)
            layer.bias_cache = np.zeros_like(layer.biases)

        # Update cache with squared current gradients
        layer.weight_cache = self.rho * layer.weight_cache + \
            (1 - self.rho) * layer.dweights**2
        layer.bias_cache = self.rho * layer.bias_cache + \
            (1 - self.rho) * layer.dbiases**2

        # Vanilla SGD parameter update + normalization
        # with square rooted cache
        layer.weights += -self.current_learning_rate * \
            layer.dweights / \
            (np.sqrt(layer.weight_cache) + self.epsilon)
        layer.biases += -self.current_learning_rate * \
            layer.dbiases / \
            (np.sqrt(layer.bias_cache) + self.epsilon)
```

```

# Call once after parameter updates
def post_update_params(self):
    self.iterations += 1

#####
#####
# Adaptive Momentum (Adam)
#####
#####
class Optimizer_Adam:
    # Initialize optimizer - set settings,
    # learning rate of 1. is default for this optimizer - changing learner rate can
    effect the ability for the
    # model to find global minimums and not get stuck in any local minimums.
    The aim is to get a low loss and a higher accuracy.
    def __init__(self, learning_rate = 0.001, decay = 0., epsilon = 1e-7,
                  beta_1 = 0.9, beta_2 = 0.999):
        self.learning_rate = learning_rate
        self.current_learning_rate = learning_rate
        self.decay = decay
        self.iterations = 0
        self.epsilon = epsilon
        self.beta_1 = beta_1
        self.beta_2 = beta_2

    # Call once before and parameter updates - this part should reduce the
    learning rate with time
    def pre_update_params(self):
        if self.decay:
            self.current_learning_rate = self.learning_rate * \
                (1. / (1. + self.decay * self.iterations))

    # Update parameters
    def update_params(self, layer):

        # If layer does not contain cache arrays,
        # create them filled with zeros
        if not hasattr(layer, 'weight_cache'):

            layer.weight_momentums = np.zeros_like(layer.weights)
            layer.weight_cache = np.zeros_like(layer.weights)
            layer.bias_momentums = np.zeros_like(layer.biases)
            layer.bias_cache = np.zeros_like(layer.biases)

        # Update momentum with current gradients
        layer.weight_momentums = self.beta_1 * \

```

```

        layer.weight_momentums + \
        (1 - self.beta_1) * layer.dweights
    layer.bias_momentums = self.beta_1 * \
        layer.bias_momentums + \
        (1 - self.beta_1) * layer.dbiases
    # Get corrected momentum
    # Self.iteration is 0 at first pass
    # and we need to start with 1 here
    weight_momentums_corrected = layer.weight_momentums /\
        (1 - self.beta_1 ** (self.iterations + 1))
    bias_momentums_corrected = layer.bias_momentums /\
        (1 - self.beta_1 ** (self.iterations + 1))
    # Update cache with squared current gradients
    layer.weight_cache = self.beta_2 * layer.weight_cache + \
        (1 - self.beta_2) * layer.dweights**2
    layer.bias_cache = self.beta_2 * layer.bias_cache + \
        (1 - self.beta_2) * layer.dbiases**2
    # Get correct cache
    weight_cache_corrected = layer.weight_cache /\
        (1 - self.beta_2 ** (self.iterations + 1))
    bias_cache_corrected = layer.bias_cache /\
        (1 - self.beta_2 ** (self.iterations + 1))

    # Vanilla SGD parameter update + normlization
    # with square roted cache
    layer.weights += -self.current_learning_rate * \
        weight_momentums_corrected /\
        (np.sqrt(weight_cache_corrected) + self.epsilon)
    layer.biases += -self.current_learning_rate * \
        bias_momentums_corrected /\
        (np.sqrt(bias_cache_corrected) + self.epsilon)

    # Call once after parameter updates
    def post_update_params(self):
        self.iterations += 1

#####
#####
# Calculating loss
#####
#####
class Loss:

    # Regularization loss calculation
    def regularization_loss(self):

```

```

# 0 by default
regularization_loss = 0

# Calculate regularization loss
# iterate all trainable layers
for layer in self.trainable_layers:

    # L1 regularization - weights
    # calculate only when factor greater than 0
    if layer.weight_regularizer_l1 > 0:
        regularization_loss += layer.weight_regularizer_l1 * \
            np.sum(np.abs(layer.weights))

    # L2 regularization - weights
    if layer.weight_regularizer_l2 > 0:
        regularization_loss += layer.weight_regularizer_l2 * \
            np.sum(layer.weights * \
                layer.weights)

    # L1 regularization - biases
    # calculate only when factor greater than 0
    if layer.bias_regularizer_l1 > 0:
        regularization_loss += layer.bias_regularizer_l1 * \
            np.sum(np.abs(layer.biases))

    # L2 regularization - biases
    if layer.bias_regularizer_l2 > 0:
        regularization_loss += layer.bias_regularizer_l2 * \
            np.sum(layer.biases * \
                layer.biases)

return regularization_loss

# Set/remember trainable layers
def remember_trainable_layers(self, trainable_layers):
    self.trainable_layers = trainable_layers

# Calculate the data and regularization losses
# given model output and ground truth values
def calculate(self, output, y, *, include_regularization = False):

    # calculate sample losses
    sample_losses = self.forward(output, y)

    # Calculate mean loss

```

```

data_loss = np.mean(sample_losses)

# Add accumulated sum of losses and sample count
self.accumulated_sum += np.sum(sample_losses)
self.accumulated_count += len(sample_losses)

# If just data loss - return it
if not include_regularization:
    return data_loss

# Return the data and regularization losses
return data_loss, self.regularization_loss()

# Calculate accumulated loss
def calculate_accumulated(self, *, include_regularization = False):

    # Calculate mean loss
    data_loss = self.accumulated_sum / self.accumulated_count

    # If just data loss - return it
    if not include_regularization:
        return data_loss

    # Return the data and regularization losses
    return data_loss, self.regularization_loss()

# Reset variables for accumulated loss
def new_pass(self):
    self.accumulated_sum = 0
    self.accumulated_count = 0

class Loss_CategoricalCrossentropy(Loss):
    # Forward pass
    def forward(self, y_pred, y_true):
        # Number of samples in a batch
        samples = len(y_pred) #want to know the total length
        # Vlip data to prevent division by 0
        # Clip both sides to not drag mean towards any value
        y_pred_clipped = np.clip(y_pred, 1e-7, 1-1e-7)

        # Probabilities for target values -
        # only if categorical labels
        if len(y_true.shape) == 1:#this means they have passed scalar values
            correct_confidences = y_pred_clipped[
                range(samples),

```

```

        y_true
    ]

    # Mask values - only for one-hot encoded labels
    elif len(y_true.shape) == 2: # this is for vectors
        correct_confidences = np.sum(
            y_pred_clipped * y_true,
            axis=1
        )

    # Losses
    negative_log_likelihoods = -np.log(correct_confidences)
    return negative_log_likelihoods

# Backward pass
def bakward(self, dvalues, y_true):

    # Number of samples
    samples = len(dvalues)
    # Number of labels in every sample
    # We'll use the first sample to count them
    labels = len(dvalues[0])

    # If lables are sparse, turn them into one-hot vector
    if len(y_true.shape) == 1:
        y_true = np.eye(labels)[y_true]

    # Calculate gradient
    self.dinputs = -y_true / dvalues
    # Normalize gradient
    self.dinputs = self.dinputs / samples

# Softmax classifier - combined Softmax activation
# and cross-entropy loss for faster backward step
class Activation_Softmax_Loss_CategoricalCrossentropy():

    # Backward pass
    def backward(self, dvalues, y_true):

        # Nuber of samples
        samples = len(dvalues)
        # If lables are one-hot encoded,
        # trun them into discrete values
        if len(y_true.shape) == 2:

```

```

        y_true = np.argmax(y_true, axis = 1)

    # Copy so we can safely modify
    self.dinputs = dvalues.copy()
    # Calculate gradients
    self.dinputs[range(samples), y_true] -= 1
    # Normalize gradient
    self.dinputs = self.dinputs / samples

# Binary cross-entropy loss
class loss_BinaryCrossentropy(Loss):

    # Forward pass
    def forward(self, y_pred, y_true):

        # Clip data to prevent division by 0
        # Clip both sides to not drag mean towards any value
        y_pred_clipped = np.clip(y_pred, 1e-7, 1 - 1e-7)

        # Calculate sample-wise loss
        sample_losses = -(y_true * np.log(y_pred_clipped) + (1 - y_true) * np.log(1
- y_pred_clipped))
        sample_losses = np.mean(sample_losses, axis = -1)

        # Return losses
        return sample_losses

    # Backward pass
    def backward(self, dvalues, y_true):

        # Number of samples
        samples = len(dvalues)
        # Number of outputs in every sample
        # We'll use the first sample to count them
        outputs = len(dvalues[0])

        # Clip data to prevent division by 0
        # Clip both sides to not drag mean towards any value
        clipped_dvalues = np.clip(dvalues, 1e-7, 1 - 1e-7)

        # Calculate gradient
        self.dinputs = -(y_true / clipped_dvalues - (1 - y_true)/(1 -
clipped_dvalues)) / outputs

        # Normalize gradient

```



```
self.dinputs = self.dinputs / samples
```

```
#####  
#####  
# calculate errors  
#####  
#####
```

```
#Mean Squared Error Loss
```

```
class Loss_MeanSquaredError(Loss): # L2 loss
```

```
    # Forward pass
```

```
    def forward(self, y_pred, y_true):
```

```
        # calculate loss
```

```
        sample_losses = np.mean((y_true - y_pred)**2, axis=-1)
```

```
        # Return losses
```

```
        return sample_losses
```

```
    # Backward pass
```

```
    def backward(self, dvalues, y_true):
```

```
        # Number of samples
```

```
        samples = len(dvalues)
```

```
        # Number of outputs in every sample
```

```
        # We'll use the first sample to count them
```

```
        outputs = len(dvalues[0])
```

```
        # Gradient on values
```

```
        self.dinputs = -2 * (y_true - dvalues)/outputs
```

```
        # Normalize gradient
```

```
        self.dinputs = self.dinputs / samples
```

```
# Mean Absolute Error loss
```

```
class Loss_MeanAbsoluteError(Loss): # L1 loss
```

```
    def forward(self, y_pred, y_true):
```

```
        # Calculate loss
```

```
        sample_losses = np.mean(np.abs(y_true - y_pred), axis = -1)
```

```
        # Return losses
```

```

    return sample_losses

# Backward pass
def backward(self, dvalues, y_true):

    #Number of samples
    samples = len(dvalues)
    # Number of outputs in every sample
    # We'll use the first sample to count them
    outputs = len(dvalues[0])

    # Calculate gradient
    self.dinputs = np.sign(y_true - dvalues)/outputs
    # Normalize gradient
    self.dinputs = self.dinputs / samples


#####
#####
# Common accuracy class
#####
#####

class Accuracy:

    # Calculates an accuracy
    # given predictions and ground truth values
    def calculate(self, predictions, y):

        # Get comparison results
        comparisons = self.compare(predictions, y)

        # Calculate and accuracy
        accuracy = np.mean(comparisons)

        # Add accumulated sum of matching values and sample count
        self.accumulated_sum += np.sum(comparisons)
        self.accumulated_count += len(comparisons)

        # Return accuracy
        return accuracy

    # Calculate accumulated accuracy
    def calculate_accumulated(self):

```

```

    # Calculate and accuracy
    accuracy = self.accumulated_sum / self.accumulated_count

    # Return the data and regularization losses
    return accuracy

# Rest variables for accumulated accuracy
def new_pass(self):
    self.accumulated_sum = 0
    self.accumulated_count = 0

# Accuracy calculation for classification model
class Accuracy_Categorical(Accuracy):

    # No initialization is needed
    def init(self, y):
        pass

    # Compares predictions to the ground truth values
    def compare(self, predictions, y):
        if len(y.shape) == 2:
            y = np.argmax(y, axis = 1)
            return predictions == y

# Accuracy calculation fro regression model
class Accuracy_Regression(Accuracy):
    def __init__(self):
        #Create precision property
        self.precision = None

    # Calculates precision value
    # based on passed in groun truth
    def init(self, y, reinit=False):
        if self.precision is None or reinit:
            self.precision = np.std(y) / 250

    # Compares predictions to the ground truth values
    def compare(self, predictions, y):
        return np.absolute(predictions - y) < self.precision

```

```
#####
#####
# Model
#####
#####
accuracyList = []
lossList = []

# Model class
class Model:

    def __init__(self):
        # create a list of network objects
        self.layers = []

        # Softmax classifier's output object
        self.softmax_classifier_output = None

    # Add objects to the model
    def add(self, layer):
        self.layers.append(layer)

    # Set loss, optimizer and accuracy
    def set(self, *, loss = None, optimizer = None, accuracy = None):

        if loss is not None:
            self.loss = loss

        if optimizer is not None:
            self.optimizer = optimizer

        if accuracy is not None:
            self.accuracy = accuracy

    # Finalize the model
    def finalize(self):

        #create and set the input layer
        self.input_layer = Layer_Input()

        # Count all the objects
        layer_count = len(self.layers)

        # Iterate the objects
        self.trainable_layers = []
```

```

# Count all the objects
for i in range(layer_count):

    # It's the first layer,
    # the previous layer object is the input layer
    if i == 0:
        self.layers[i].prev = self.input_layer
        self.layers[i].next = self.layers[i+1]

    # All layers except for the first and the last
    elif i < layer_count - 1:
        self.layers[i].prev = self.layers[i-1]
        self.layers[i].next = self.layers[i+1]

    # The last layer - the next object is the loss
    # Also let's save aside the reference to the last object
    # whose output is the model's output
    else:
        self.layers[i].prev = self.layers[i-1]
        self.layers[i].next = self.loss
        self.output_layer_activation = self.layers[i]

    # If layer contains an attribute called "weights",
    # it's a trainable layer -
    # add it to the list of trainable layers
    # we don't need to check for biases -
    # checking for weights is enough
    if hasattr(self.layers[i], 'weights'):
        self.trainable_layers.append(self.layers[i])

# Update loss object with trainable layers
if self.loss is not None:

    self.loss.remember_trainable_layers(
        self.trainable_layers
    )

# If output activation is Softmax and
# loss function is Categorical Cross-Entropy
# create an object of combined activation
# and loss function containing
# faster gradient calculation
if isinstance(self.layers[-1], Activation_Softmax) and \
    isinstance(self.loss, Loss_CategoricalCrossentropy):
    # Create an object of combined activation

```

```

# and loss function
self.softmax_classifier_output = \
    Activation_Softmax_Loss_CategoricalCrossentropy()

# Train the model
def train(self, X, y, *, epochs=1, batch_size=None,
        print_every=1, validation_data=None):

    # Initialize accuracy object
    self.accuracy.init(y)

    # Default value if batch size is not set
    train_steps = 1

    # If there is validation data passed,
    # set default number of steps for validation as well
    if validation_data is not None:
        validation_steps = 1

    # For better readability
    X_val, y_val = validation_data

    # Calculate number of steps
    if batch_size is not None:
        train_steps = len(X) // batch_size
        # Dividing rounds down. If there are some remaining
        # data, but not a full batch, this won't include it
        # Add 1 to include this not full batch
        if train_steps * batch_size < len(X):
            train_steps += 1

    if validation_data is not None:
        validation_steps = len(X_val) // batch_size
        # Dividing rounds down. If there are some remaining
        # data, but not a full batch, this won't include it
        # Add 1 to include this not full batch
        if validation_steps * batch_size < len(X_val):
            validation_steps += 1

    # Main training loop
    for epoch in range(1, epochs + 1):

        # print epoch number
        print(f'epoch: {epoch}')

```

```

# Reset accumulated values in loss and accuracy objects
self.loss.new_pass()
self.accuracy.new_pass()

# Iterate over steps
for step in range(train_steps):
    # If batch size is not set
    # train using one step and full dataset
    if batch_size is None:
        batch_X = X
        batch_y = y

    # Otherwise slice a batch
    else:
        batch_X = X[step*batch_size:(step+1)*batch_size]
        batch_y = y[step*batch_size:(step+1)*batch_size]

    # Perform the forward pass
    output = self.forward(batch_X, training = True)

    # Calculate loss
    data_loss, regularization_loss = \
        self.loss.calculate(output, batch_y,
                            include_regularization = True)
    loss = data_loss + regularization_loss

    # Get predictions and calculate and accuracy
    predictions = self.output_layer_activation.predictions(output)
    accuracy = self.accuracy.calculate(predictions, batch_y)

    # Perform backward pass
    self.backward(output, batch_y)

    # Optimize (update parameters)
    self.optimizer.pre_update_params()
    for layer in self.trainable_layers:
        self.optimizer.update_params(layer)
    self.optimizer.post_update_params()

    # Print a summary
    if not step % print_every or step == train_steps - 1:

        # Print out summary
        print(f'step: {step}, ' +
              f'acc: {accuracy:.2f}, ' +
              f'loss: {loss:.3f}, (' +

```

```

        f'data_loss: {data_loss:.3f}, ' +
        f'reg_loss: {regularization_loss:.3f}), ' +
        f'lr: {self.optimizer.current_learning_rate:.4f}')

    # Get and print epoch loss and accuracy
    epoch_data_loss, epoch_regularization_loss = \
        self.loss.calculate_accumulated(include_regularization = True)
    epoch_loss = epoch_data_loss + epoch_regularization_loss
    epoch_accuracy = self.accuracy.calculate_accumulated()

    print(f'training, ' +
          f'acc: {accuracy:.2f}, ' +
          f'loss: {loss:.3f}, (' +
          f'data_loss: {data_loss:.3f}, ' +
          f'reg_loss: {regularization_loss:.3f}), ' +
          f'lr: {self.optimizer.current_learning_rate:.4f}')

    # If there is the validation data
    if validation_data is not None:

        # Evaluate the model - this accounts for the first validation and loss
        # output we see. the second one is from the evaluation at the bottom.
        self.evaluate(*validation_data,
                      batch_size=batch_size)

    # Performs forward pass
    def forward(self, X, training):

        # Call forward method on the input layer
        # this will set the output property that
        # the first layer in "prev" object is expecting
        self.input_layer.forward(X, training)

        # Call forward method of every object in a chain
        # Pass output of the previous object as a parameter
        # for layer in self.layers:
        layer.forward(layer.prev.output, training)

        # "layer" is now the last object from the list,
        # Return its output
        return layer.output

    # Perform a backward pass
    def backward(self, output, y):
        # If softmax_classifier

```



```

if self.softmax_classifier_output is not None:
    # First call backward method
    # on the combined activation/loss
    # this will set diinputs property
    self.softmax_classifier_output.backward(output, y)

    # Since we'll not call backward method of the last layer
    # which is Softmax activation
    # as we used combined activation/loss
    # object, let's set dinputs in this object
    self.layers[-1].dinputs = \
        self.softmax_classifier_output.dinputs

    # Class backward method going through
    # all the objects but last
    # in reversed order passing dinputs as a parameter
    for layer in reversed(self.layers[:-1]):
        layer.backward(layer.next.dinputs)

    return

# First call backward method on the loss
# this will set dinputs property that the last
# layer will try to access shortly
self.loss.backward(output, y)

# Call backward method going through all the objects
# in reversed order passing dinputs as a parameter
for layer in reversed(self.layers):
    layer.backward(layer.next.dinputs)

# Evaluates the model using passed in dataset
def evaluate(self, X_val, y_val, *, batch_size=None):

    # Default value if batch size is not being set
    validation_steps = 1

    # Calculate number of steps
    if batch_size is not None:
        validation_steps = len(X_val) // batch_size
        # Dividing rounds down. If there are some remaining
        # data but not a full batch, this won't include it
        # Add '1' to include this not full batch
        if validation_steps * batch_size < len(X_val):
            validation_steps += 1

```

```

# Reset accumulated values in loss and accuracy objects
self.loss.new_pass()
self.accuracy.new_pass()

# Iterate over steps
for step in range(validation_steps):

    # If batch size is not set -
    # train using one step and full dataset
    if batch_size is None:

        batch_X = X_val
        batch_y = y_val

    # Otherwise slice a batch
    else:
        batch_X = X_val[step*batch_size:(step+1)*batch_size]
        batch_y = y_val[step*batch_size:(step+1)*batch_size]

    # Perform the forward pass
    output = self.forward(batch_X, training = False)

    # Calculate the loss
    self.loss.calculate(output, batch_y)

    # Get predictions and calculate and accuracy
    predictions = self.output_layer_activation.predictions(output)
    self.accuracy.calculate(predictions, batch_y)

# Get and print validation loss and accuracy
validation_loss = self.loss.calculate_accumulated()
validation_accuracy = self.accuracy.calculate_accumulated()

# Print a summary
print(f'validation, ' +
      f'acc: {validation_accuracy:.3f}, ' +
      f'loss: {validation_loss:.3f}' +
      '\n')

# Retrieves and returns parameters of trainable layers
def get_parameters(self):

    # Create a list for parameters
    parameters = []

```

```

# Iterable trainable layers and get their parameters
for layer in self.trainable_layers:
    parameters.append(layer.get_parameters())

# Return a list
return parameters

# Update the model with new parameters
def set_parameters(self, parameters):

    # Iterate over the parameters and layers
    # and update each layers with each set of the parameters
    for parameter_set, layer in zip(parameters, self.trainable_layers):
        layer.set_parameters(*parameter_set)

# Saves the parameters to a file
def save_parameters(self, path):

    # Open a file in the binary-write mode
    # and save parameters to it
    with open(path, 'wb') as f:
        pickle.dump(self.get_parameters(), f)

# Loads the weights and updates a model instance with them
def load_parameters(self, path):

    # Open file in the binary-read mode,
    # Load weights and update trainable layers
    with open(path, 'rb') as f:
        self.set_parameters(pickle.load(f))

# Saves the model
def save(self, path):

    # Make a deep copy of current model instance
    model = copy.deepcopy(self)

    # Reset accumulated values in loss and accuracy objects
    self.loss.new_pass()
    self.accuracy.new_pass()

    # Remove data from the input layer
    # and gradients from the loss object
    model.input_layer.__dict__.pop('output', None)
    model.loss.__dict__.pop('dinputs', None)

```

```

# For each layer remove inputs, output and dinputs properties
for layer in model.layers:
    for property in ['inputs', 'output', 'dinputs', 'dweights', 'dbiases']:
        layer.__dict__.pop(property, None)

# Open a file in the binary-write mode and save the model
with open(path, 'wb') as f:
    pickle.dump(model, f)

# Loads and returns a model
@staticmethod
def load(path):

    # Open file in the binary-read mode, load a model
    with open(path, 'rb') as f:
        model = pickle.load(f)

    # Return a model
    return model

# Predicts on the samples
def predict(self, X, *, batch_size=None):

    # Default value is batch size is not being set
    prediction_steps = 1

    # Calculate number of steps
    if batch_size is not None:
        prediction_steps = len(X) // batch_size
        # Dividing rounds down. If there are some remaining
        # data, but not a full batch, this won't include it
        # Add '1' to include this not full batch
        if prediction_steps * batch_size < len(X):
            prediction_steps += 1

    # Model outputs
    output = []

    # Iterate over steps
    for step in range(prediction_steps):

        # If batch size is not set -
        # train using one step and full dataset
        if batch_size is None:
            batch_X = X

```

```

        # Otherwise slice a batch
        else:
            batch_X = X[step*batch_size:(step+1)*batch_size]

        # Perform the forward pass
        batch_output = self.forward(batch_X, training = False)

        # Append batch prediction to the list of predictions
        output.append(batch_output)

    # Stack and return results
    return np.vstack(output)

#
#
#
#
#
#
#
#
#####
#####
# Loading Data
#####
#####

# Loads a MNIST dataset
def load_mnist_dataset(dataset, path):

    # Scan all the directories and create a list of lables
    print('Scanning all directories and creating a list of lables')
    labels = os.listdir(os.path.join(path, dataset))

    #labels.sort()
    #print(labels[1:])

    # Create lists for samples and lables
    X = []
    y = []

    # For each lable folder
    print('Reading the image')
    # I had to use labels[1:] because there is an invisible file
    # .DS_Store in the folder that kept creating NotADirectoryError: [Errno 20]
    for label in labels[1:]:

```

```

# And for each image in given folder
for file in os.listdir(os.path.join(
    path, dataset, label
)):

    # Read the image
    image = cv2.imread(os.path.join(
        path, dataset, label, file
    ), cv2.IMREAD_UNCHANGED)

    # And append it and a label to the lists
    X.append(image)
    y.append(label)

# Convert the data to proper numpy arrays and return
return np.array(X), np.array(y).astype('uint8')

# MNIST dataset (train + test)
def create_data_mnist(path):

    # Load both sets separately
    print('Loading training data')
    X, y = load_mnist_dataset('train', path)
    print('loading test data')
    X_test, y_test = load_mnist_dataset('test', path)

    # And return all the data
    return X, y, X_test, y_test

#####
#####
#
#
#
#
#
#
#
#####
#####
# Input data
#####
#####

# we can load our data by doing

```

```
X, y, X_test, y_test = create_data_mnist('fashion_mnist_images')
```

```
#####  
#####
```

```
# Shuffle the training dataset  
print('Shuffling the training dataset')  
keys = np.array(range(X.shape[0]))  
np.random.shuffle(keys)  
X = X[keys]  
y = y[keys]
```

```
#####  
#####
```

```
# Scale features
```

```
# scale images to be between the range of -1 and 1 by taking each pixel value,  
# subtracting half the maximum of all pixel values (i.e.,  $255/2 = 127.5$ ).  
# NB: We could also scale our data between 0 and 1 by simply dividing it by 255  
(the maximum value).  
print('Scaling both training and test data')  
X = (X.reshape(X.shape[0], -1).astype(np.float32) - 127.5)/127.5  
X_test = (X_test.reshape(X_test.shape[0], -1).astype(np.float32) - 127.5)/127.5
```

```
#####  
#####
```

```
#  
#  
#  
#  
#  
#  
#  
#
```

```
#####  
#####
```

```
# Artificial Neural Layers
```

```
# Model Object
```

```
# 1x512 densely-connected neural network (2 hidden layers with 512 neurons)
```

```
#####  
#####
```

```

# Instantiate the model
model = Model()

# add layers
model.add(layer_Dense(X.shape[1], 128)) # input layer
model.add(Activation_ReLU())
model.add(Layer_Dropout(0.1)) # dropout layer
model.add(layer_Dense(128,128)) # 1st hidden layer
model.add(Activation_ReLU())
model.add(Layer_Dropout(0.1)) # dropout layer
model.add(layer_Dense(128,10)) # 2nd hidden layer
model.add(Activation_Softmax()) # output layer

#print(model.layers)

# Set loss, optimizer and accuracy objects
model.set(
    loss=Loss_CategoricalCrossentropy(), #loss=Loss_MeanSquaredError(),
    optimizer=Optimizer_Adam(decay=1e-3),
    accuracy=Accuracy_Categorical() #accuracy=Accuracy_Regression()
)

# Finalize the model
model.finalize()

# Train the model
model.train(X, y, validation_data = (X_test, y_test),
            epochs=10, batch_size=128, print_every=100)

# Retrive and print parameters
parameters = model.get_parameters()

# New model

# Instance the model
model = Model()

# add layers
model.add(layer_Dense(X.shape[1], 128)) # input layer
model.add(Activation_ReLU())
model.add(Layer_Dropout(0.1)) # dropout layer
model.add(layer_Dense(128,128)) # 1st hidden layer
model.add(Activation_ReLU())

```



```

model.add(Layer_Dropout(0.1)) # dropout layer
model.add(layer_Dense(128,10)) # 2nd hidden layer
model.add(Activation_Softmax()) # output layer


# Set loss and accuracy objects
model.set(
    loss=Loss_CategoricalCrossentropy(), #loss=Loss_MeanSquaredError(),
    accuracy=Accuracy_Categorical() #accuracy=Accuracy_Regression()
)


# Finalize the model
model.finalize()


# Set model with parameters instead of training it
model.set_parameters(parameters)


#####
#####
# Saving and loading the parameters


# Save paraters (weights and biasies)
#model.save_parameters('fashion_mnist.parms')


# Load saved parameters (weights and biasies)
#model.load_parameters('fashion_mnist.parms')


#####
#####
# Saving and loading the model


# Save model
#model.save('fashion_mnist.model')


#####
#####
#####
# Present the name of the prediction


# I will need to change this according to the parkour move that I want the
system to recognise
fashion_mnist_labels = {
    0: 'T-shirt/top',

```

```
1: 'Trouser',
2: 'Pullover',
3: 'Dress',
4: 'Coat',
5: 'Sandal',
6: 'Shirt',
7: 'Sneaker',
8: 'Bag',
9: 'Ankle boot'
}
```

```
#####
#####
# Get image for prediction
```

```
# get image data and change it to grey scale
image_data = cv2.imread('prediction_images/pants.png',
cv2.IMREAD_GRAYSCALE)
```

```
# Resize the plot so that it is the same size as the test data images
image_data = cv2.resize(image_data, (28, 28))
```

```
# invert the pixels so that they look like the images in the test data i.e. black
background with white clothing.
image_data = 255 - image_data
```

```
# Reshape and scale pixel data
image_data = (image_data.reshape(1, -1).astype(np.float32)) - 127.5 /127.5
```

```
#####
#####
# Load model
model = Model.load('fashion_mnist.model')
```

```
# Evaluate the model
#model.evaluate(X_test, y_test)
```

```
#####
#####
# Confidences
```

```
# Predict on the first 5 samples from validation dataset
#confidences = model.predict(X_test[:5])
```

```
# Predict on the image
confidences = model.predict(image_data)
```

```
# Print the confidence result
#print('confidences: ', confidences)
```

```
#####
#####
# Predictions
```

```
predictions = model.output_layer_activation.predictions(confidences)
#print('predictions: ', predictions) # will show only the numbers so we need to
convert the class number back to the name of what it is
```

```
# Get label name from label index
prediction = fashion_mnist_labels[predictions[0]]

print(prediction)
```